

Attention

Tarushii Goel

April 2020

1 Overview

Recurrent Neural Networks (RNNs) provide outstanding results on a variety of Natural Language Processing tasks; however, they continue to struggle in capturing long-term dependencies in complex sentences. Several advancements, such as LSTMs & GRUs were designed to tackle the problem of learning long-term dependencies, but perhaps one of the most influential is Attention. This lecture will begin by discussing the origin of Attention as a improvement to Seq2Seq for Neural Machine Translation (NMT), discuss and compare a variety of attention mechanisms, and conclude by exploring self-attention multi-head attention, which serve as the groundwork for recent advancements in Transformers and BERT.

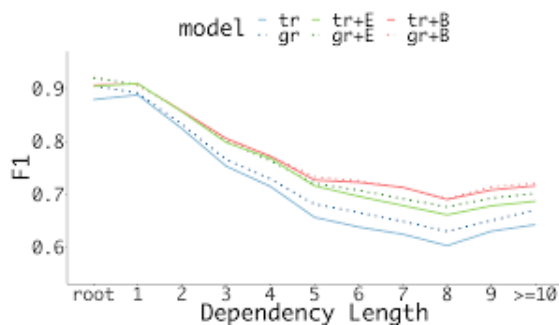


Figure 1: The Long-Term Dependency Problem (Graph of BERT variants, which are state-of-the-art)

2 Recap

To refresh your memory, I'll briefly review RNNs and Encoder-Decoders models. RNNs are recurrent; in each iteration they are fed the next input in a sequence and the previous hidden state, which carries accumulated information from other parts of the sentence, and produce a new hidden state.




Figure 2: A Diagram of an RNN

Seq2Seq are a unique RNN architecture used for NMT. It is composed to two RNNs; an encoder, which 'encodes' the sentences into a 'context' vector (the last encoder hidden state, a latent representation of the entire sentence), and a decoder, which takes the 'context' vector and decodes it into a translated sentence.




Figure 3: A Diagram of an RNN

3 What is Attention?

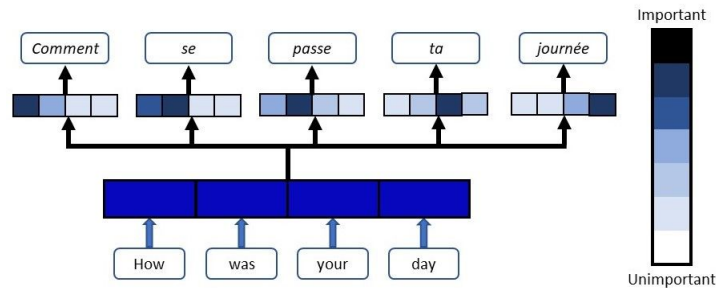
When we think about the English word “Attention”, we know that it means directing your focus at something and taking greater notice. The Attention mechanism in Deep Learning is based off this concept of directing your focus, and it pays greater attention to certain factors when processing the data.

In broad terms, Attention is one component of a network’s architecture, and is in charge of managing and quantifying the interdependence:

1. Between the input and output elements (General Attention)
2. Within the input elements (Self-Attention)

Say we have the sentence “How was your day”, which we would like to translate to the French version - “Comment se passe ta journée”. What the Attention component of the network will do for each word in the output sentence is map the important and relevant words from the input sentence and assign higher weights to these words, enhancing the accuracy of the output prediction.

Now, the question remains, how do we calculate attention weights?



4 Types of Attention Calculation

General Attention is used to improve Seq2Seq models. Traditionally, the only information carried over from the encoder to the decoder is the last encoder hidden state, but with Attention, a weighted sum of the encoder hidden states of every time-step is appended to the input. Given a query q and a set of key-

Figure 5: General Attention

values. The alignment model, is a row of weights applied to the decoder hidden state and the encoder hidden states. To calculate an 'alignment score' for any, calculate the dot-product $s_i^T * h_i$. Once the 'alignment scores are calculated for $h_0 - n$, apply a softmax to get the weights.

4.2 Self-Attention

With self-attention, each hidden state attends to the previous hidden states of the same RNN. Here s_t is the query while the decoder hidden states s_0 to $s_t - 1$ represent both the keys and the values.




Figure 6: Self-Attention

4.3 Multi-Head Attention

When we have multiple queries q , we can stack them in a matrix Q . Each query can be thought of as a different 'head.' In the figure below, two attention heads are displayed (beige and green). The query word is "it". The first attention head is focusing more on the "animal" while the second in green is focusing on "tired". These multi-head mechanisms apparently help in translation tasks.




Figure 7: Multi-Head Attention

5 Code

I coded a standard Seq2Seq model with attention in PyTorch and put the code below. I highly recommend implementing the techniques I discussed in this lecture on your own and using this as a reference. Let me know if you have any questions!

```

1 #data is from European Parliament proceedings, download here: http
  ://www.statmt.org/europarl/
2 PATH = 'YOUR_PATH' #set to the path to your data
3 en_data = open(PATH + 'europarl-v7.fr-en.en', encoding='utf-8').
  read().split('\n')
4 fr_data = open(PATH + 'europarl-v7.fr-en.fr', encoding='utf-8').
  read().split('\n')
5
6 import spacy
7 import torch
8 import torchtext
9 from torch import nn
10 import torch.nn.functional as F
11
12 device = torch.device("cuda" if torch.cuda.is_available() else "cpu
  ")
13
14 en_field = torchtext.data.Field(sequential=True, use_vocab=True,
  init_token='<sos>', eos_token='<eos>', tokenize='spacy',
  tokenizer_language='en')
15 fr_field = torchtext.data.Field(sequential=True, use_vocab=True,
  init_token = '<sos>', eos_token = '<eos>', tokenize='spacy',
  tokenizer_language='fr')
16
17 #Now, covert the data to a csv file to take advantage of torchtext'
  s versatile TabularDataset class
18 import pandas as pd
19 from sklearn.model_selection import train_test_split
20

```

```

21 raw_data = {'English': [line for line in en_data], 'French': [line
    for line in fr_data]}
22 df = pd.DataFrame(raw_data, columns=["English", "French"])
23 # create train and validation set
24 train, val = train_test_split(df, test_size=0.1)
25 train.to_csv(PATH + "train.csv", index=False)
26 val.to_csv(PATH + "val.csv", index=False)
27
28 train_set, val_set = torchtext.data.TabularDataset.splits(path=PATH,
    train='train.csv', validation='val.csv', format='csv', fields
    =[('English', en_field), ('French', fr_field)])
29
30 #loads 100-dimensional pre-trained glove word embeddings
31 en_field.build_vocab(train_set, val_set, vectors='glove.6B.100d')
32 fr_field.build_vocab(train_set, val_set, vectors='glove.6B.100d')
33
34 train_iter, val_iter = torchtext.data.BucketIterator.splits((
    train_set, val_set), batch_size=32, sort_key=lambda x: len(x.
    French), shuffle=True)
35
36 class Embedder(nn.Module):
37     def __init__(self, embedding):
38         super().__init__()
39         self.embed = nn.Embedding(embedding.shape[0], embedding.
    shape[1])
40         self.embed.weight.data.copy_(embedding)
41         self.embed.weight.requires_grad = False
42     def forward(self, input_sequence):
43         return self.embed(input_sequence)
44
45 class Embedder(nn.Module):
46     def __init__(self, embedding):
47         super().__init__()
48         self.embed = nn.Embedding(embedding.shape[0], embedding.
    shape[1])
49         self.embed.weight.data.copy_(embedding)
50         self.embed.weight.requires_grad = False
51     def forward(self, input_sequence):
52         return self.embed(input_sequence)
53
54 class Encoder(nn.Module):
55     def __init__(self, hidden_size, embedding, num_layers=1,
    dropout=0.0):
56         super(Encoder, self).__init__()
57         self.hidden_size = hidden_size
58         self.embedding = Embedder(embedding)
59         self.gru = nn.GRU(embedding.shape[1], hidden_size,
    num_layers=num_layers, dropout=dropout, bidirectional=True,
    batch_first=True)
60     def forward(self, input_sequence):
61         embedded = self.embedding(input_sequence)
62         # x = nn.utils.rnn.pack_padded_sequence(x, lens) # unpad
63         outputs, hidden_state = self.gru(embedded) # gru returns
    hidden state of all timesteps as well as hidden state at last
    timestep
64         # pad the sequence to the max length in the batch
65         # output, _ = nn.utils.rnn.pad_packed_sequence(output)

```

```

66         # The output of a GRU has shape (seq_len, batch, hidden_size
        * num_directions)
67         # Because the Encoder is bidirectional, combine the results
        from the
68         # forward and reversed sequence by simply adding them
        together.
69         outputs = outputs[:, :, :self.hidden_size] + outputs[:, :,
        self.hidden_size:]
70         return outputs, hidden_state
71
72     class Decoder(nn.Module):
73     def __init__(self, batch_size, hidden_size, embedding,
74                 num_layers=1, drop_prob=0.1):
75         super(Decoder, self).__init__()
76         self.batch_size = batch_size
77         self.embedding_size = embedding.shape[1]
78         self.embedding = Embedder(embedding)
79         self.attn = nn.Linear(hidden_size * 2, 1)
80         self.gru = nn.GRU(embedding.shape[1]+hidden_size,
81                             hidden_size, num_layers=num_layers, dropout=drop_prob,
82                             batch_first=True)
83         self.classifier = nn.Linear(hidden_size, embedding.shape
84                                     [0])
85
86     def forward(self, decoder_hidden, encoder_outputs, inputs):
87         # Embed input words
88         embedded = self.embedding(inputs)
89         #Assumed size of decoder_hidden -> (num_layers, batch_size,
90         embedding_size) size of encoder_outputs -> (batch_size,
91         sentence_len, embedding_size)
92         #need to convert length of decoder_hidden to (batch_size,
93         sentence_len, embedding_size)
94         self.sequence_len = encoder_outputs.shape[1]
95         decoder_hidden = torch.sum(decoder_hidden, axis=0)
96         attn_inp = decoder_hidden.unsqueeze(1).repeat(1, self.
97         sequence_len,1)
98         weights = self.attn(torch.cat((attn_inp, encoder_outputs),
99         dim = 2)).squeeze()
100         normalized_weights = F.softmax(weights)
101         attn_applied = torch.bmm(normalized_weights.unsqueeze(1),
102         encoder_outputs)
103         cat_input = torch.cat((embedded, attn_applied), axis=2)
104         output, hidden_state = self.gru(cat_input, decoder_hidden.
105         unsqueeze(0))
106         output = self.classifier(output).squeeze()
107         return output, hidden_state
108
109     def train(encoder, decoder, encoder_opt, decoder_opt, criterion,
110             input, target):
111         #set both to train mode
112         encoder.train()
113         decoder.train()
114         #pass through encoder
115         enc_output, enc_hidden = encoder(target)
116         #initialize input to '<sos>' tokens and decoder hidden state to
117         final encoder hidden state
118         dec_input, dec_hidden = target[:, 0].unsqueeze(1), enc_hidden

```

```

106     loss = 0
107     for i in range(1, target.shape[1]):
108         dec_input, dec_hidden = decoder(dec_hidden, enc_output,
109                                         dec_input)
110         loss += criterion(dec_input, target[:, i])
111         topv, topi = dec_input.topk(1)
112         dec_input = topi.detach() # detach from history as input
113     loss.backward()
114     encoder_opt.step()
115     decoder_opt.step()
116     return loss
117
118 hidden_size = 200
119 batch_size = 2
120 epochs = 1
121 encoder = Encoder(hidden_size, en_field.vocab.vectors).to(device)
122 decoder = Decoder(batch_size, hidden_size, fr_field.vocab.vectors).
123             to(device)
124 encoder_opt = torch.optim.Adam([param for param in encoder.
125                                 parameters() if param.requires_grad == True], lr=1.0e-4)
126 decoder_opt = torch.optim.Adam([param for param in decoder.
127                                 parameters() if param.requires_grad == True], lr=1.0e-4)
128 criterion = nn.CrossEntropyLoss(ignore_index=1)
129
130 for i in range(epochs):
131     losses = []
132     print('Epoch %x:' % i, end='')
133     c = 0
134     for batch in train_iter:
135         input = batch.English.t().to(device)
136         target = batch.French.t().to(device)
137         if (torch.max(input) > 116730):
138             continue
139         losses.append(train(encoder, decoder, encoder_opt,
140                             decoder_opt, criterion, input, target).item())
141         if (c % 10 == 0):
142             print('.', end='')
143         c+=1
144     print(' loss: ', sum(losses)/len(losses))

```

6 References

I got a lot of the explanation and graphics for Attention from this blog post:
<https://blog.floydhub.com/attention-mechanism/>.